

Computer Network I

Reti di Calcolatori I

Università di Napoli Federico II – Scuola Politecnica e delle Scienze di Base
Corso di Laurea in Informatica

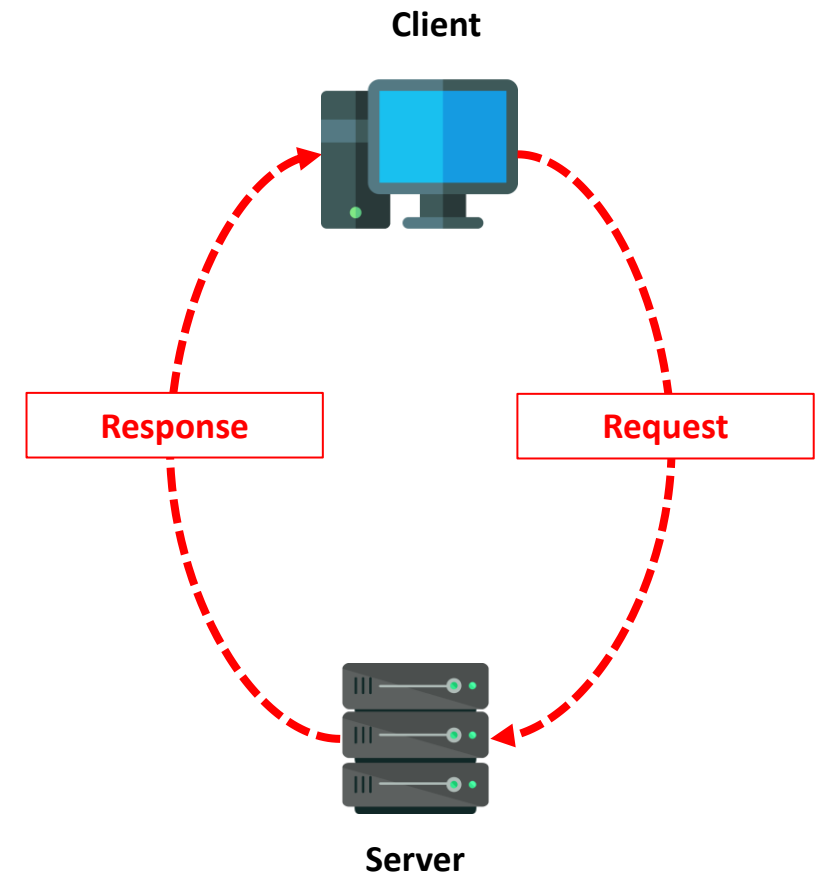
Riccardo Caccavale
(riccardo.caccavale@unina.it)



Application Layer

HTTP: Message Format


- In HTTP protocol we have 2 different formats for the request and the response messages.
- **Request:** specifies the command (method) that the HTTP server has to perform (e.g., give me a web page, fill a specific form, etc.).
- **Response:** reports the outcome of the command (success, fail, etc.) and possible data (e.g., the requested file).
- Both messages are written in ordinary ASCII text, so they are easily readable by humans.



Application Layer

HTTP: Request Message Format

- The request message format include the following fields:
 - The **method**: specifies the requested command to be executed by the server.
 - The **URL**: is used to identify the object on which we want to operate.
 - The **version**: specifies the HTTP version (e.g., HTTP/1.1).
 - The **header lines**: contain the parameters of the request, the number and the type of these lines are not fixed. Each line include the **name** and the **value** of the parameter.
 - For instance, we can here specify if we want **persistent** or **non-persistent** connection.
 - Custom headers can also be used.
 - The **body**: is method-specific and contains data that are potentially associated with the command (e.g., text used to fill a form).




The fields are separated by special characters:

- The **sp** is space character.
- The **cr** is carriage return (**\r**).
- The **lf** is line feed (**\n**).

Application Layer

HTTP: Request Message Format (Methods)

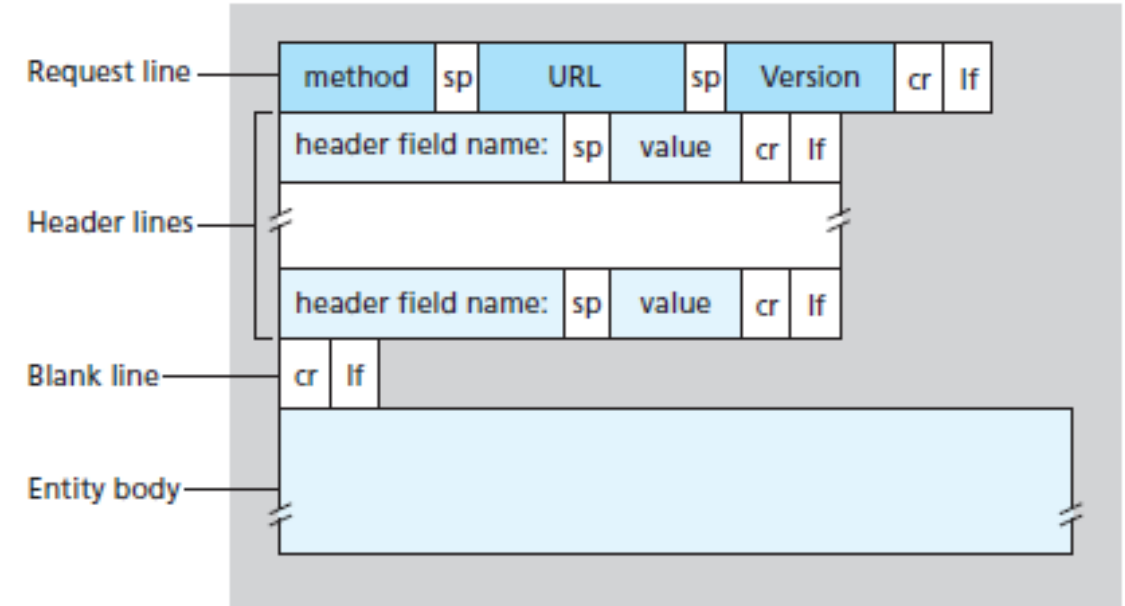
- The **GET** method is used to retrieve objects (resources) from the server.
 - This is one of the most used commands as every time we browse a new webpage the associated file must be retrieved.
 - In the GET method the **body is empty**.
- The **POST** method is used to set info inside objects (resources) of the server.
 - A request generated with a form does not necessarily use the POST method, HTML forms often use the GET method and include the inputted data (in the form fields) in the requested URL.
 - The body contains the info to be posted.



Application Layer

HTTP: Request Message Format (Methods)

- The **HEAD** method is similar to the GET method, when a server receives a request with the HEAD method, it responds with an HTTP message, but the object is not returned.
 - Application developers often use the HEAD method for debugging.
- The **PUT** method is often used in conjunction with Web publishing tools.
 - It allows a user to upload an object to a specific path (directory) on a specific Web server.
 - The PUT method is also used by applications that need to upload objects to Web servers.
- The **DELETE** method allows a user, or an application, to delete an object on a Web server.



Application Layer

HTTP: Request Message Format (Example)

- In this example we have a **browser** (Mozilla/5.0), implementing “HTTP/1.1”, that is **requesting** the object “/somedir/page.html” to the “www.someschool.edu” web server.
- We have **five lines**:
 - The **GET request line** (resource and version).
 - The **four header lines** (host, connection, user-agent, accept-language)...

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
Connection: close
User-agent: Mozilla/5.0
Accept-language: fr
```

Application Layer

HTTP: Request Message Format (Example)

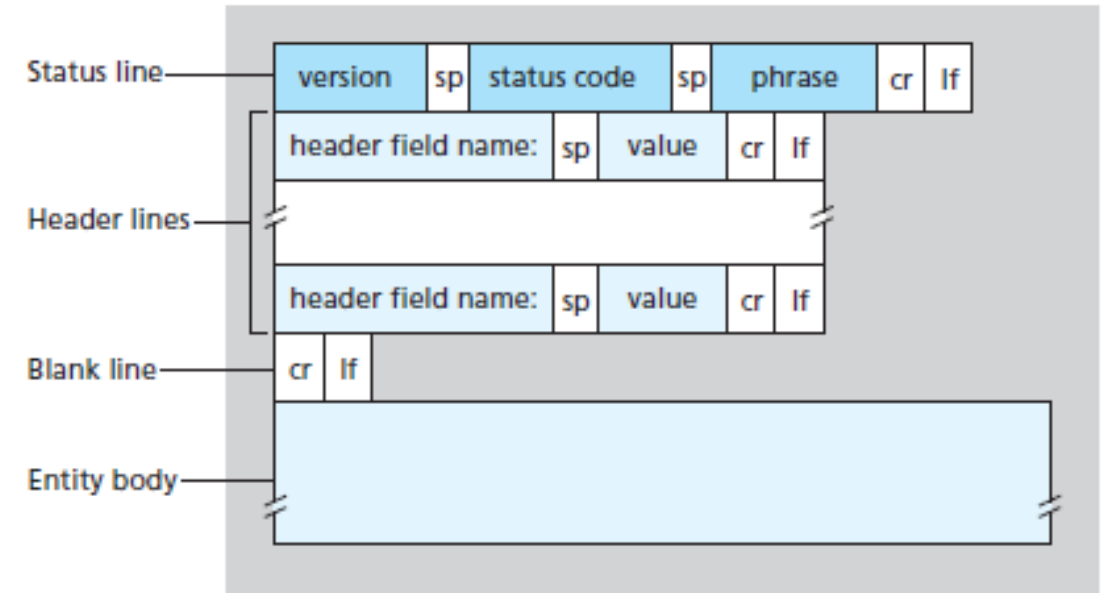
- The line “**Host: www.someschool.edu**” specifies the **host on which the object resides**.
 - **Destination is not always the next host**: a message could be **forwarded by another host** (e.g., proxy server), so the address in the TCP is different from the real target host.
- The line “**Connection: close**” specifies that the browser is asking for a non-persistent connection (i.e., close after finish).
- The line “**User-agent: Mozilla/5.0**” specifies the **browser’s type**.
 - This is sometimes useful because the **server can send different versions** of the same object (each version is addressed by the same URL) **depending on the type**.
- The line “**Accept-language: fr**” indicates that the **user prefers to receive a French version** of the object (if exists).

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
Connection: close
User-agent: Mozilla/5.0
Accept-language: fr
```

Application Layer

HTTP: Response Message Format


- The general format of the response message is similar to the request one.
- In this case, instead of a request line, we have a **status line** that reports the outcome of the command that includes:
 - The **version**: reports the HTTP version of the server's response.
 - The **status code**: a code (number) that specifies the outcome of the command.
 - The **phrase**: contains the result of the request.



Application Layer

HTTP: Response Message Format (Status Codes)

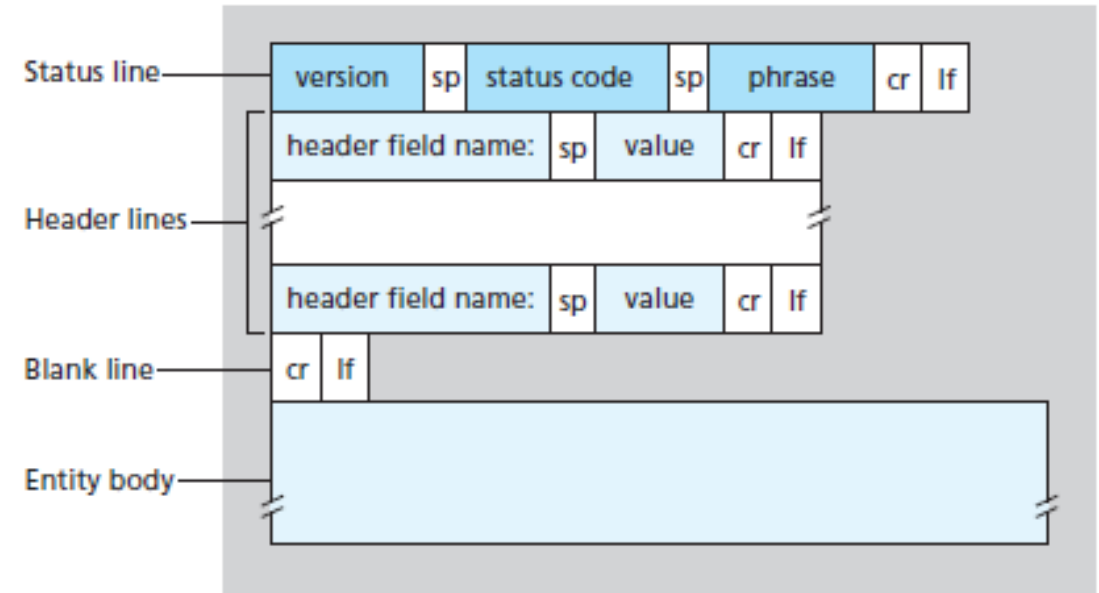
- **Status codes** are divided into classes:
 - 100-199 **informational**: info regarding the request.
 - 200-299 **successful**: the request has been executed successfully.
 - 300-399 **redirection**: there are additional actions needed from the client.
 - 400-499 **client-error**: the request cannot be executed because of a client issue.
 - 500-599 **server-error**: the request cannot be executed because of a server issue.



Application Layer

HTTP: Response Message Format (Status Codes)

- Some typical examples are:
 - **200 OK**: request succeeded, and the information is returned in the response.
 - **301 Moved Permanently**: requested object has been permanently moved; the new URL is specified in the “Location” header of the response message.
 - **400 Bad Request**: generic error code indicating that the request could not be understood by the server.
 - **404 Not Found**: The requested document does not exist on this server.
 - **505 HTTP Version Not Supported**: The requested HTTP protocol version is not supported by the server.



Application Layer

HTTP: Response Message Format (Example)

- In this example, the status line indicates that **the server is using HTTP/1.1 and that everything is OK** (the server has found and is sending the requested object).
- There are **five header lines** (connection, date, server, last-modified, content-length, content-type).
- The **body contains the requested object** (represented by data data data data data ...).

```
HTTP/1.1 200 OK
Connection: close
Date: Tue, 18 Aug 2015 15:44:04 GMT
Server: Apache/2.2.3 (CentOS)
Last-Modified: Tue, 18 Aug 2015 15:11:03 GMT
Content-Length: 6821
Content-Type: text/html

(data data data data data ...)
```

Application Layer

HTTP: Response Message Format (Example)

- The line “**Connection: close**” informs the client that the **connection will be closed** after this message (non-persistent).
- The line “**Date: ...**” indicates the time and **date when the HTTP response was created** and sent by the server.
- The line “**Server: ...**” indicates that the message was **generated by an Apache Web server** (similar to user-agent).
- The line “**Last-Modified: ...**” indicates the time and date when the **object was created or last modified**.
 - Useful in case of caching, as cached files can be outdated.
- The line “**Content-Length: ...**” indicates the **number of bytes in the object**.
- The line “**Content-Type: ...**” indicates that the **object is HTML text** (here, the object type is officially indicated by this header line and not by the file extension).

```
HTTP/1.1 200 OK
Connection: close
Date: Tue, 18 Aug 2015 15:44:04 GMT
Server: Apache/2.2.3 (CentOS)
Last-Modified: Tue, 18 Aug 2015 15:11:03 GMT
Content-Length: 6821
Content-Type: text/html

(data data data data data ...)
```

Application Layer

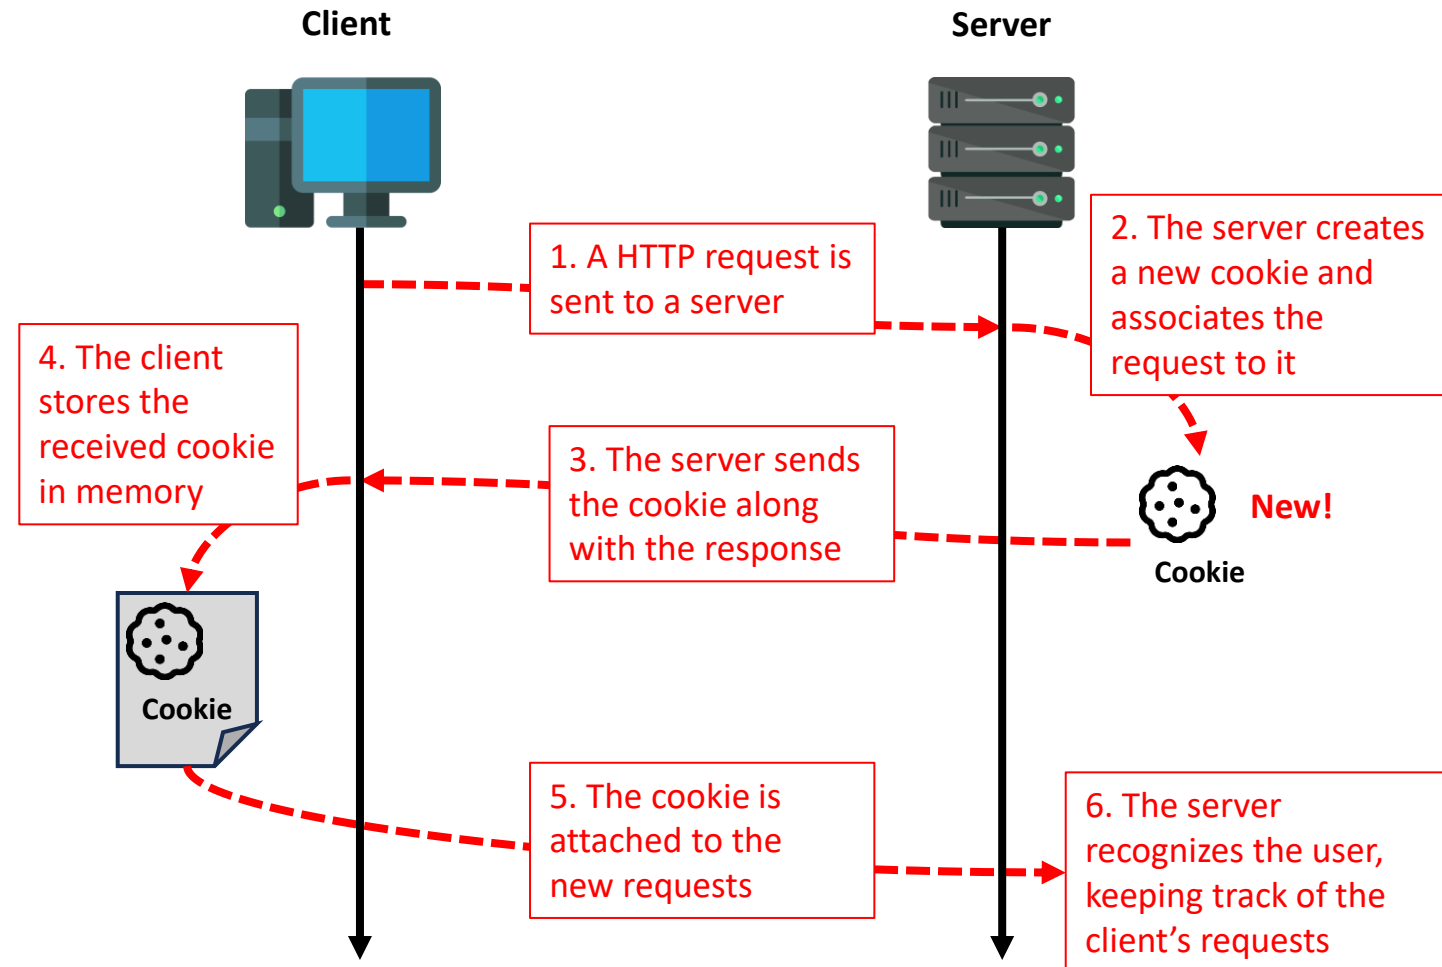
HTTP: Cookies

- An HTTP **server is typically stateless**, this simplifies server design reduces the use of resources and allow servers to handle thousands of simultaneous TCP connections.
- Pure statelessness is also a strong limitation as **several web functions are client-specific** (for example, Amazon's cart depends on the client, Netflix suggests contents based on client's preferences, etc.).
- For these purposes, HTTP uses cookies. A **cookie** is a digital token (alphanumeric ID) used by servers to identify a specific client.
 - Cookies allow web sites to keep track of users, most of the major commercial Web sites use cookies.
 - Cookies may also have attributes (e.g., expiration date).

Application Layer

HTTP: Cookies


- A **cookie** is created by the **server** and delivered to the client.
- Cookie technology involves four main components:
 1. A **“Set-cookie” header** line in the HTTP response message.
 2. A **“Cookie” header** line in the HTTP request message.
 3. A **file** on the client system (managed by the user’s browser).
 4. A **back-end database** on the server.



Application Layer

HTTP: Cookies (Example)

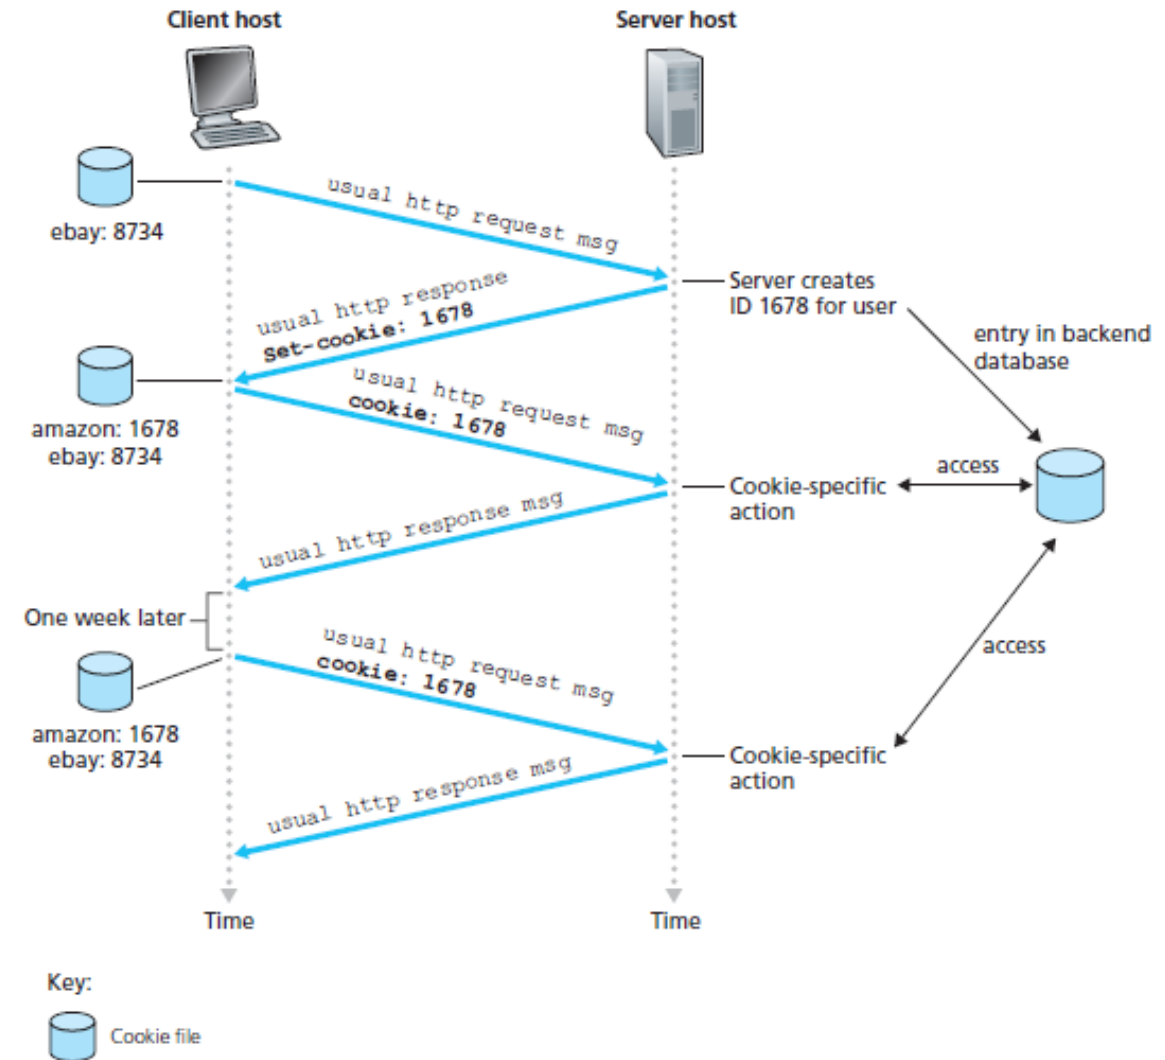
- Let's assume a client host, which uses eBay regularly, to **contact Amazon.com for the first time**.
- The client already has a cookie for eBay, but it **has no cookie for Amazon**.
- When the request comes into the Amazon Web server, **the server creates a cookie** (ID number) and an associated entry in its back-end database.
- The Amazon Web server then responds to client's browser and **includes in the HTTP response a "Set-cookie" header line**, which contains the ID (Set-cookie: 1678).



Application Layer

HTTP: Cookies (Example)

- When the response is received, **the browser appends a line to a special cookie file** including:
 - the hostname of the server.
 - the ID number of the cookie.
- From now on, the requests from the client to Amazon **will be associated to the new cookie** by including the header line `Cookie: 1678` to the HTTP messages.
- Amazon server is then able to **track the client's activity through the database**. It knows exactly which pages user 1678 visited, in which order, and at what times.
- If the client returns to Amazon's site **one week later**, the browser will continue to put the header line `Cookie: 1678` in the request messages.



Application Layer


HTTP: Cookies

- Amazon (and other web sites) uses cookies to provide different services:
 - **Shopping cart service**, the server can maintain a list of intended purchases during browsing.
 - **Products recommendations**, based on the visited web pages.
 - **User's registration**, by associating user's info to the cookie in the database (credit-card, name, e-mail, address) so you don't have to re-insert them every time.
- Although cookies often simplify the Internet shopping experience for the user, they are **controversial** because they can also be considered as an invasion of privacy.
- Using a combination of **cookies** and user-supplied **account information**, a Web site can learn a lot about a user and potentially sell this information to a third party.

Application Layer

HTTP: Web Caching


- A **Web cache** (also called a **proxy server**) is a network entity that satisfies HTTP requests on the behalf of an origin Web server.
- The Web cache has its own disk storage and **keeps copies of recently requested objects** in this storage.
- A browser can be configured so that all the HTTP **requests are first directed to the Web cache** to check if a copy of the requested object is available.



Application Layer

HTTP: Web Caching


- Let's assume a browser to request the object `http://www.someschool.edu/campus.gif` passing through a web cache:
 - The **browser establishes a TCP connection** to the Web cache and sends an HTTP request for the object to the Web cache.
 - The Web cache **checks if it has a copy of the object** stored locally. If so, the Web cache returns the object within an HTTP response message to the client browser.
 - If the **Web cache does not have the object**, the Web cache opens a TCP connection to the origin server (`www.someschool.edu`) and sends an HTTP request for the object.
 - When the **Web cache receives the object**, it **stores a copy in its local storage** and sends a copy, within an HTTP response message, to the client browser.



Application Layer

HTTP: Web Caching


- Note that a cache is both a server (when providing objects) and a client (when requesting objects) at the same time.
- Web caching has seen deployment in Internet for two reasons:
 1. A Web cache can substantially **reduce the response time for a client request**, particularly if a high-speed connection stands between the client and the cache.
 2. Web caches can substantially **reduce traffic** of a company or institution toward Internet, in so reducing costs due to bandwidth.



Application Layer

HTTP: Web Caching


- Web caches are **often installed into a company/institution local network** to speed-up and reduce traffic.
- A **hit** happens when a cache successfully provide an object without contacting the original server.
- The hit rate, i.e., the fraction of requests that are satisfied by a cache, typically **ranges from 0.2 to 0.7**.
 - Hit rate increases when more clients use the cache.
- This means that up to 70% of requests can be served locally.



Application Layer

HTTP: Web Caching

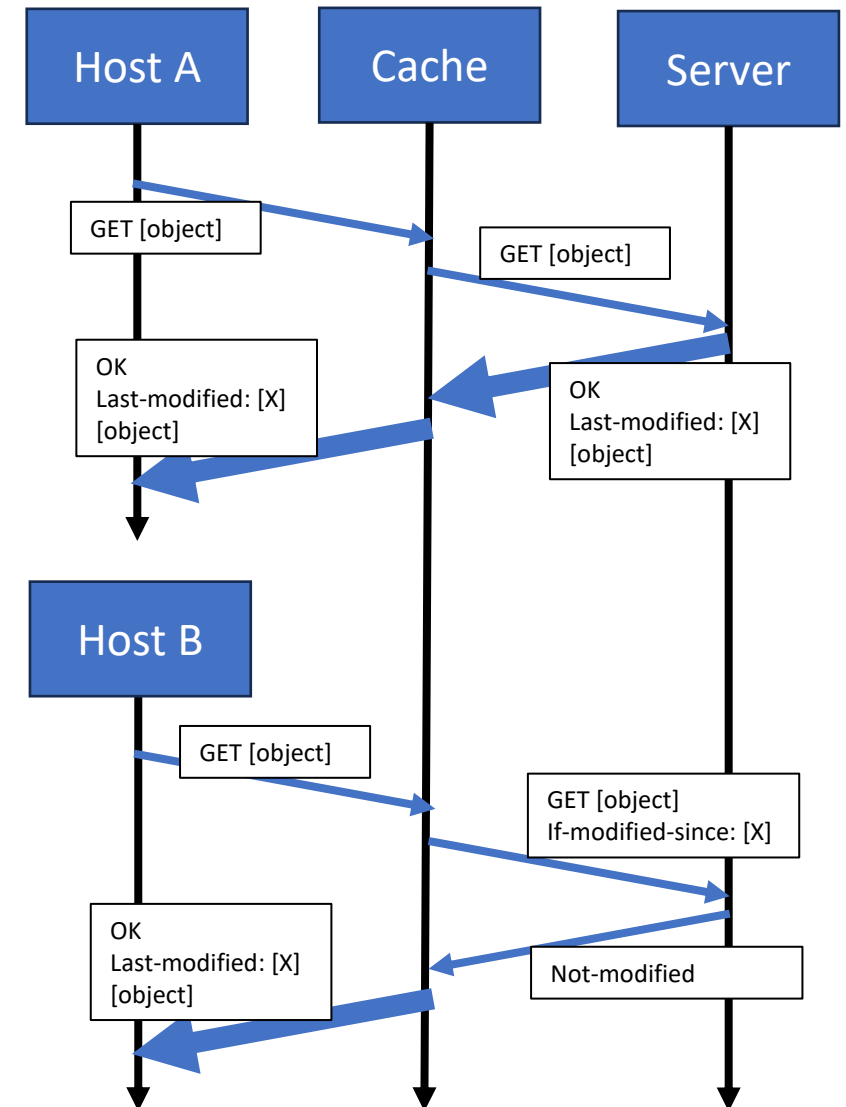
- The UNINA network also provides an institutional proxy (proxy.unina.it).
- Besides performance, **proxy servers can also be used to access to institutional services.**
- As requests are forwarded by the institutional proxy, from the original server's standpoint all **requests are coming from the institution.**
- There are also several commercial/free **proxy servers available on Internet.**



Application Layer

HTTP: Web Caching

- Web caching introduces a new problem: the copy of an object residing in the **cache may be outdated**.
- To avoid this issue, HTTP has a mechanism that allows a cache to **verify the stored object**.
- The **conditional GET** is an HTTP request message including a GET method and “If-Modified-Since:” header.
- The cache **checks if the object is up to date** and, if so, the stored version is sent back to the host (no further communication needed).



Application Layer

HTTP: Web Caching

- Caching **is also performed locally** by browsers.
- The principle is the same as servers, the browser stores objects locally so they no need to be retrieved from the server.
- This is a common techniques in modern browsers as **it drastically improves performance**.
- The local version of **the object may be not updated**, generating errors (quite frequent).

