

# Computer Network I

Reti di Calcolatori I

Università di Napoli Federico II – Scuola Politecnica e delle Scienze di Base  
Corso di Laurea in Informatica

Riccardo Caccavale  
([riccardo.caccavale@unina.it](mailto:riccardo.caccavale@unina.it))



# Application Layer

## Creating Network Applications

- So far, we have looked at several network applications and protocols, we will now see **how network applications can be created**.
- Most network applications include a **client-side** program and a **server-side** program that communicate through the network.
- When these two programs are executed, a client process and a server process are created, and these processes communicate with each other by reading from, and writing to, **sockets**.
- When creating a network application, the developer's main task is therefore to write the code for both the client and server programs.

# Application Layer

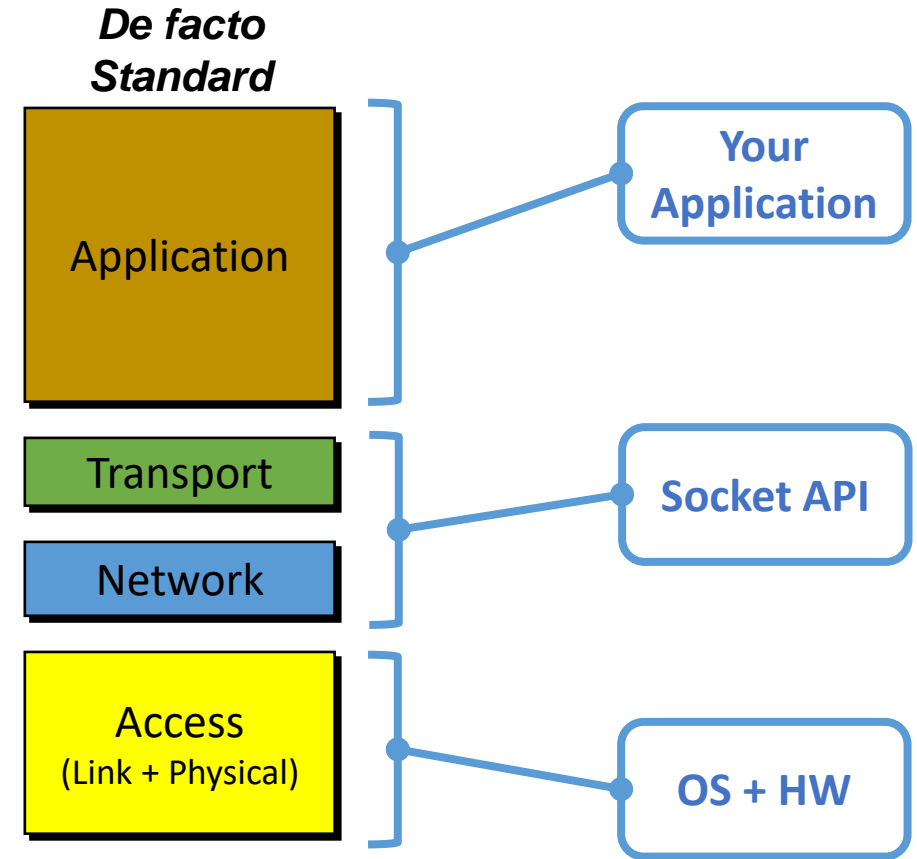
## Creating Network Applications

- There are two types of network applications:
  - Applications relying on a **standard protocol**. There are several “open” protocols whose rules are known. In this case, the client and server programs must conform to the rules to be compatible with the protocol.
  - Applications relying on a **proprietary protocol**. In this case the client and server programs employ a custom protocol (not been openly published in an RFC or elsewhere). In this case, a single developer (or development team) creates both the client and server programs.
- There are two transport protocols that can be used:
  - **TCP** (Transmission Control Protocol), which is **connection oriented** and provides a reliable byte-stream channel through which data flows between two end systems.
  - **UDP** (User Datagram Protocol), which is **connectionless** and sends independent packets of data from one end system to the other, without any guarantees about delivery.

# Application Layer

## Creating Network Applications

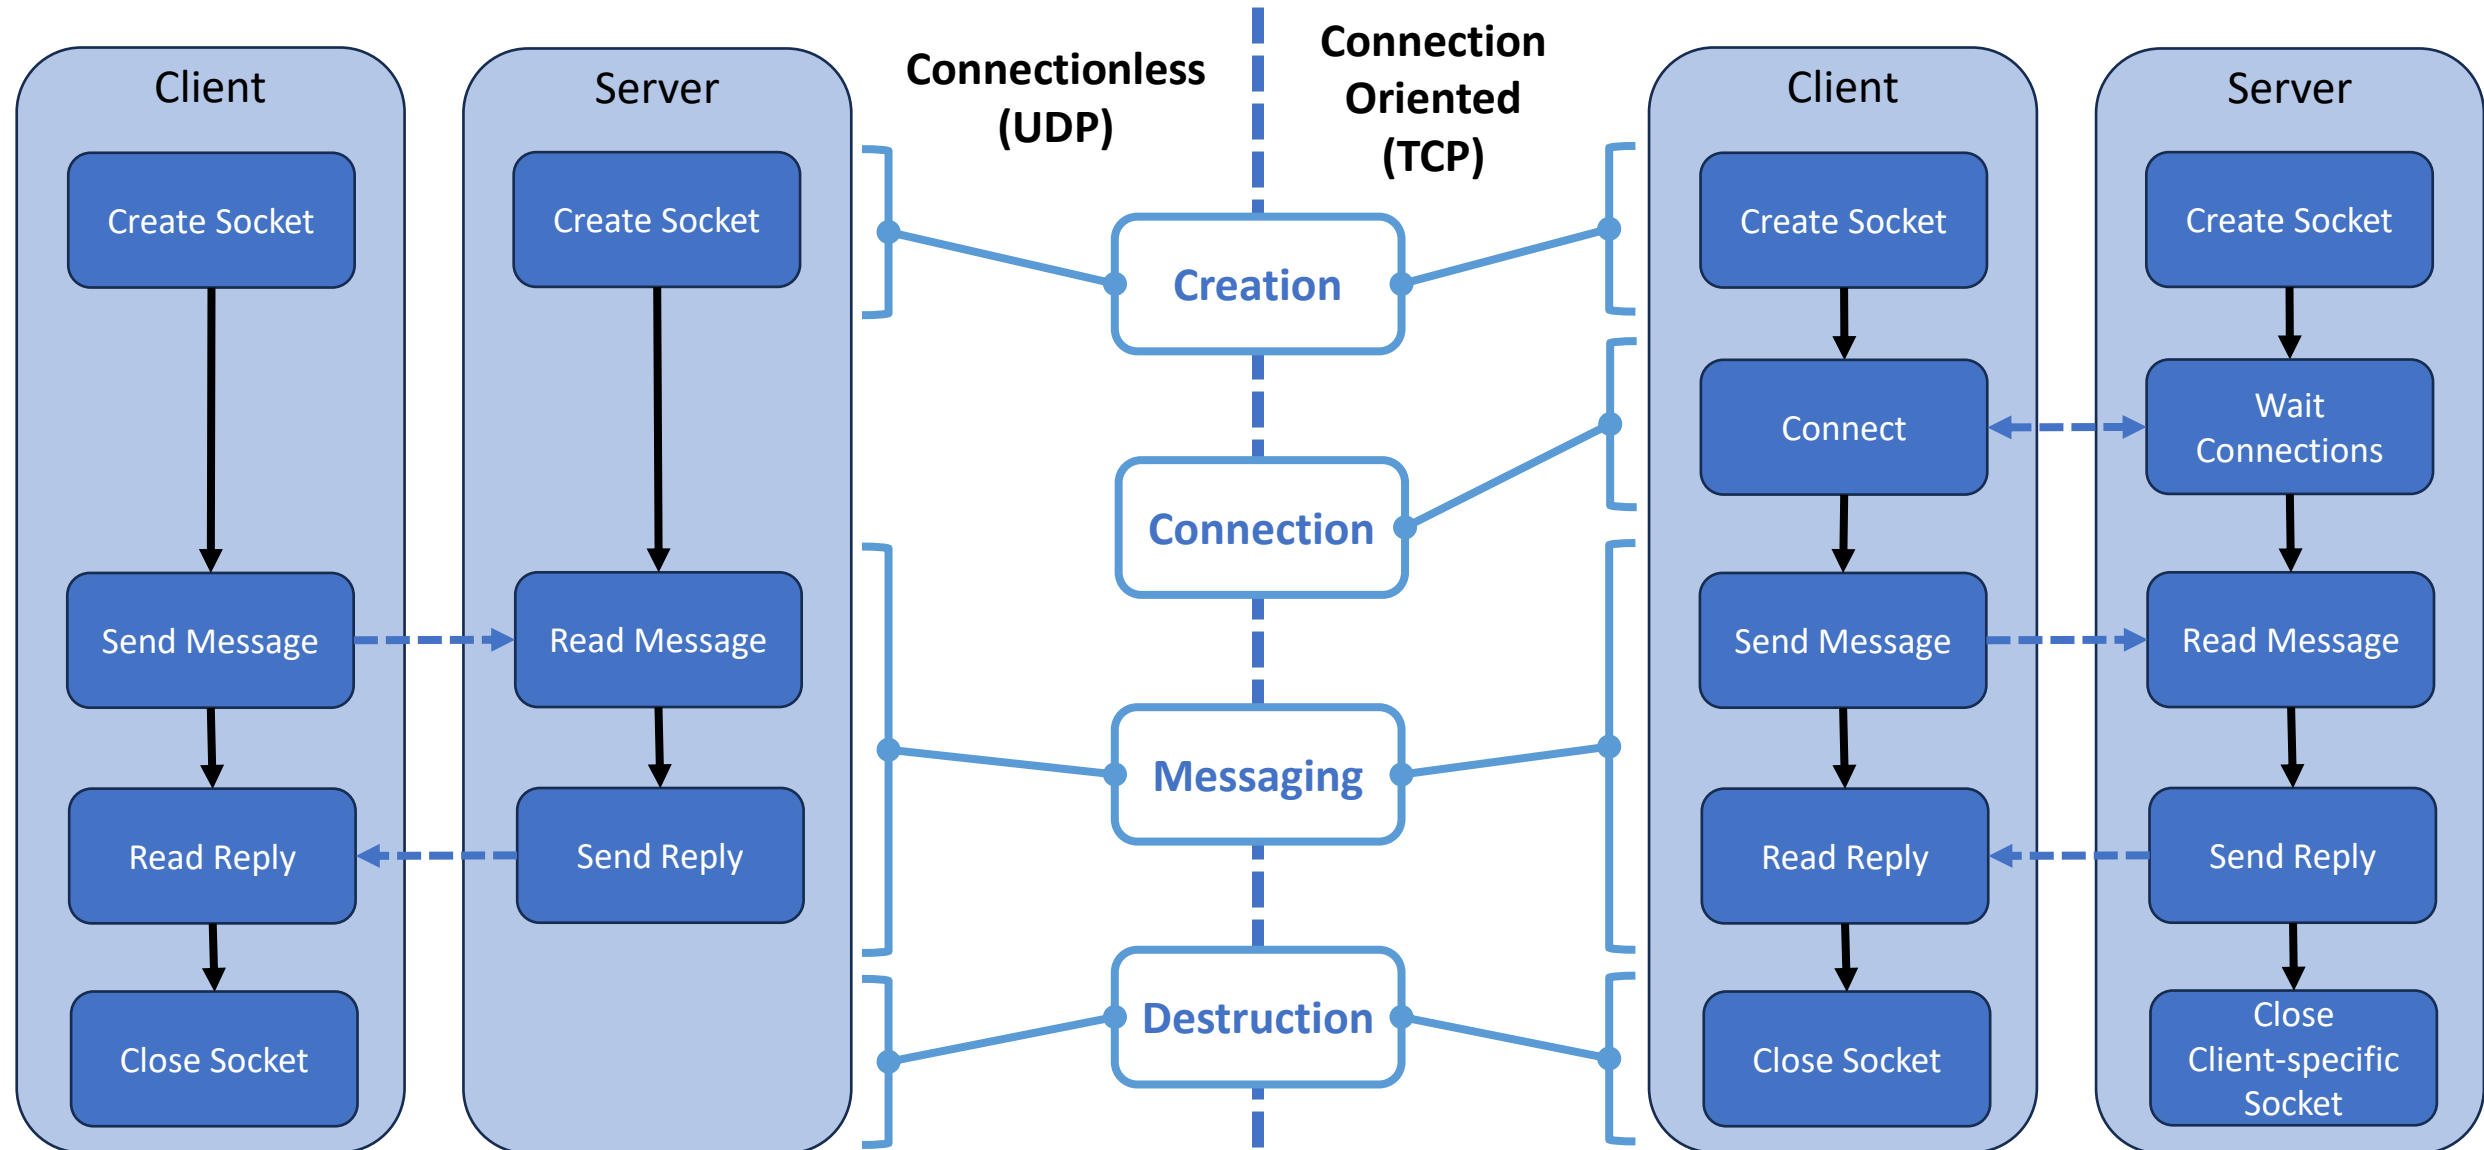
- **Sockets** are a central element for the creation of network applications.
- They typically manages (black-box) the transport (UDP/TCP) and the network (IP) layers functionalities.
- Note that there are several applications (middleware) that may **wrap sockets** simplifying their creation and providing more complex functionalities.
- However, **basic API are still widely used.**



Where layers are implemented

# Application Layer

## Connection-oriented vs. Connectionless



# Application Layer

## UDP and TCP

- UDP and TCP sockets implement the previous procedure to allow connectionless and connection-oriented communication, respectively.
- The basic APIs providing UDP and TCP sockets are typically available in almost **all programming languages and operating systems**:
  - C, C++, C#, Java, Perl, Python, Matlab, etc.
- The underlying implementation and the usage can be more or less different depending on the configuration of the machines.
- In Unix domain we use **Berkeley sockets** (aka BSD or POSIX sockets) for both TCP and UDP connections. These are native C APIs.

# Application Layer

## Sockets: Definitions

- Internet-domain sockets in C rely on some structure to define the addresses and ports of the sending/receiving hosts:

```
#include <netinet/in.h>
#include <arpa/inet.h>      // for inet_addr()
#include <sys/types.h>      // some custom types are defined here

struct sockaddr_in {
    short      sin_family;   // family of the address, typically set to AF_INET (IPv4)
    unsigned short sin_port; // port number, e.g. htons(3490)
    struct in_addr sin_addr; // see struct in_addr below
    char        sin_zero[8]; // typically zeros, now this structure has same size as sockaddr and can be casted to it
};

struct in_addr {
    unsigned long s_addr; // IP address, can be loaded with inet_addr() or set to INADDR_ANY for localhost
};
```

# Application Layer

## Sockets: Creation

- Sockets are created in C by using the `socket()` function:

```
#include <sys/socket.h>

int sockfd = socket(int domain, int type, int protocol);
```

### Where:

- domain: specifies the family as in previous structure (`AF_INET`).
- type: specifies the type of socket to be created.
  - `SOCK_STREAM` -> TCP socket
  - `SOCK_DGRAM` -> UDP socket
- protocol: Specifies a particular protocol to be used within the socket.
  - This is often 0, so no specific protocol is used.
- sockfd: is the file descriptor that identifies the newly created socket.
  - Notice that sockets also follow the Unix philosophy that “everything is a file” where I/O sources are treated as files (and associated to a file descriptor).



# Application Layer

## Sockets: Binding

- We can associate a socket to local address and port by using the bind() function:

```
#include <sys/socket.h>

int val = bind(int socket, const struct sockaddr *address, socklen_t address_len);
```

### Where:

- socket: is the file descriptor of the socket we want to bind.
  - address: pointer to a sockaddr structure specifying on which port and address to bind.
    - address.sin\_addr often set to INADDR\_ANY so connections from all addresses are listened (multiple IPs).
  - address\_len: specifies the length of the sockaddr structure pointed to by the address argument (we can use sizeof() to get it).
  - val: return value which is 0 on success, and -1 otherwise.
- 
- Notice that this function **is used on servers** as we must bind the socket to a specific port but is **optional on clients** (the OS assigns a free port automatically).

# Application Layer

## Sockets: Send and Receive (UDP)

- Processes of sending or receiving messages in UDP are performed by the `sendto()` and `recvfrom()` functions:

```
#include <sys/socket.h>
```

```
int ob = sendto(int osock, const void *obuf, size_t olen, int oflags, const struct sockaddr *oaddr, socklen_t oaddr_len);
```

```
#include <sys/socket.h>
```

```
int ib = recvfrom(int isock, void *ibuf, size_t ilen, int iflags, struct sockaddr *iaddr, socklen_t *iaddr_len);
```

### Where:

- `osock/isock`: are the file descriptors on which to send/receive a message.
- `obuf/ibuf`: are pointers to the buffers containing a message to send/receive.
- `olen/ilen`: are lengths of the messages in bytes.
- `oflags/iflags`: specifies flags (typical values are 0 or `MSG_WAITALL` to wait until all `olen/ilen` bytes are sent/received).
- `oaddr/iaddr`: pointers to `sockaddr` structures containing the receiving/sending address.
- `oaddr_len/iaddr_len`: length in bytes of the `sockaddr` structure.
- `ob/ib`: number of bytes that are actually sent/received.

# Application Layer

## Sockets: Closing

- Sockets can be closed by using the close() function:

```
#include <unistd.h>

int val = close(int socket);
```

Where:

- socket: is the file descriptor of the socket to close
- val: return value which is 0 on success, and -1 otherwise.

# Application Layer

## UDP Socket Programming (C/C++)

// client-side (includes omitted)

```
int main() {  
    int sockfd;  
    char buffer[1024];  
    const char *hello = "Hello from client";  
    struct sockaddr_in servaddr;
```

```
    if ( ( sockfd = socket(AF_INET, SOCK_DGRAM, 0)) < 0 ) {  
        perror("socket creation failed");  
        exit(EXIT_FAILURE);  
    }
```

```
    memset(&servaddr, 0, sizeof(servaddr));
```

```
    servaddr.sin_family = AF_INET;  
    servaddr.sin_port = htons(8080);  
    servaddr.sin_addr.s_addr = INADDR_ANY;
```

```
    int n;  
    socklen_t len;  
  
    sendto(sockfd, (const char *)hello, strlen(hello),  
        MSG_CONFIRM, (const struct sockaddr *) &servaddr, sizeof(servaddr));  
    std::cout<<"Hello message sent."<<std::endl;
```

```
    n = recvfrom(sockfd, (char *)buffer, 1024,  
        MSG_WAITALL, (struct sockaddr *) &servaddr, &len);  
    buffer[n] = '\0';  
    std::cout<<"Received \""<<buffer<<"\"<<std::endl;
```

```
    close(sockfd);
```

```
    return 0;  
}
```

**Note:** to talk to another computer use `inet_addr()` instead of `INADDR_ANY` (e.g., `inet_addr("192.168.1.10")` )

### UDP Client

Create UDP  
Socket

Send Message

Read Reply

Close Socket

### UDP Server

Create UDP  
Socket

Read Message

Send Reply

//server-side (includes omitted)

```
int main() {  
    int sockfd;  
    char buffer[1024];  
    const char *hello = "Hello from server";  
    struct sockaddr_in servaddr, cliaddr;
```

```
    if ( ( sockfd = socket(AF_INET, SOCK_DGRAM, 0)) < 0 ) {  
        perror("socket creation failed");  
        exit(EXIT_FAILURE);  
    }
```

```
    memset(&servaddr, 0, sizeof(servaddr));  
    memset(&cliaddr, 0, sizeof(cliaddr));
```

```
    servaddr.sin_family = AF_INET;  
    servaddr.sin_addr.s_addr = INADDR_ANY;  
    servaddr.sin_port = htons(8080);
```

```
    if ( bind(sockfd, (const struct sockaddr *)&servaddr, sizeof(servaddr)) < 0 ) {  
        perror("bind failed");  
        exit(EXIT_FAILURE);  
    }
```

```
    socklen_t len = sizeof(cliaddr);  
    int n;
```

```
    n = recvfrom(sockfd, (char *)buffer, 1024,  
        MSG_WAITALL, (struct sockaddr *) &cliaddr, &len);  
    buffer[n] = '\0';  
    std::cout<<"Received \""<<buffer<<"\"<<std::endl;
```

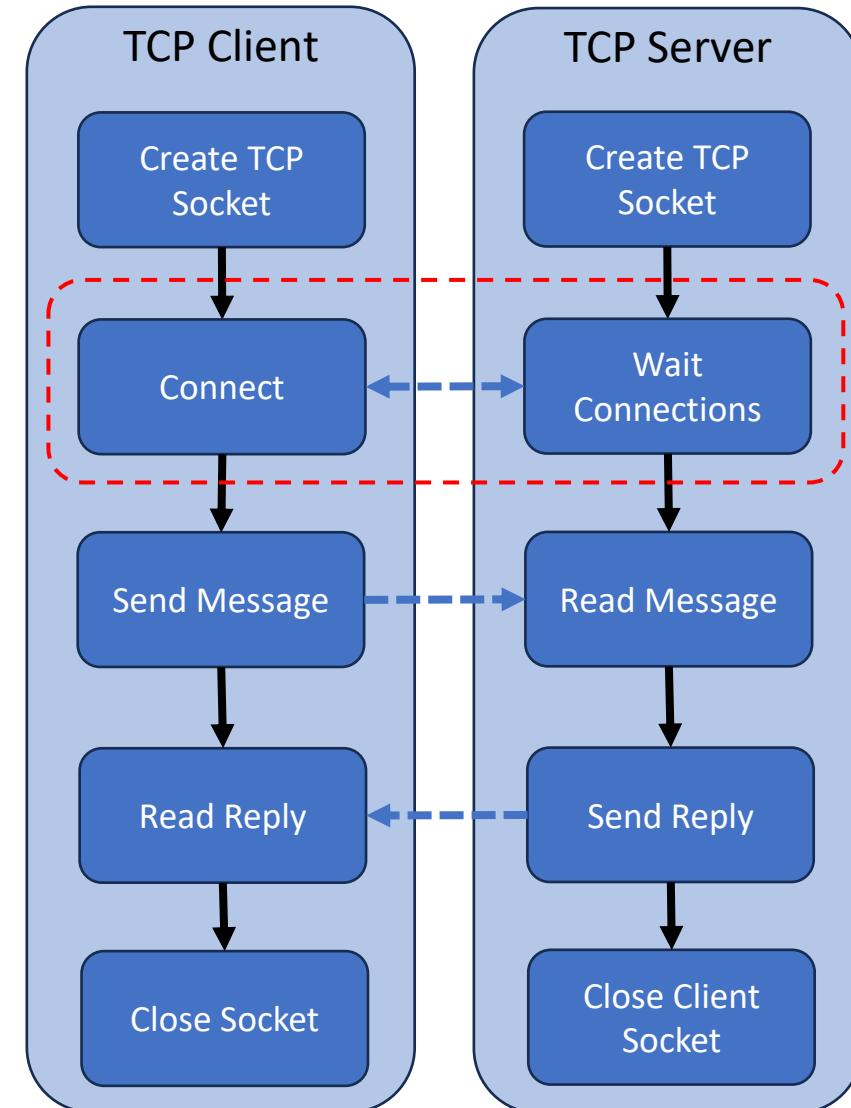
```
    sendto(sockfd, (const char *)hello, strlen(hello),  
        MSG_CONFIRM, (const struct sockaddr *) &cliaddr, len);  
    std::cout<<"Hello message sent."<<std::endl;
```

```
    close(sockfd);  
    return 0;  
}
```

# Application Layer

## From UDP to TCP

- Unlike UDP, in TCP client and server need to **handshake** before messages can be transmitted.
- We use **two sockets on the server-side**:
  - A **welcoming socket**, which is always on, and allows clients to perform handshakes with the server.
  - A **client-specific socket**, which is created after the handshake, and is used by client and server to communicate.
- The three-way handshake, which takes place within the transport layer, is completely **invisible to the client and the server programs**.
- Once the connection is accepted by the server (handshake success) **the communication moves to the newly created socket**.



# Application Layer

## Sockets: Connection (TCP only)

- Client-side TCP connection is performed by using the connect() function:

```
#include <sys/socket.h>
```

```
int val = connect(int socket, const struct sockaddr *address, socklen_t address_len);
```

### Where:

- socket: is the file descriptor of the socket we use to connect.
- address: pointer to a sockaddr structure containing the server address.
- address\_len: specifies the length of the sockaddr structure pointed to by the address argument (we can use sizeof() to get it).
- val: return value which is 0 on success, and -1 otherwise.


# Application Layer

## Sockets: Wait-for-connection (TCP only)

- Server-side TCP wait-for-connection is performed through two functions: `listen()` and `accept()`, which works in combination with the clients `connect()`.

- The `listen()` opens a queue where incoming connections are stored.

- The `accept()` opens the client-specific socket.



# Application Layer

## Sockets: Wait-for-connection (TCP only)

- Server-side TCP wait-for-connection is performed by the `listen()` and `accept()` functions:

```
#include <sys/socket.h>

int val = listen(int socket, int backlog);
int new_sockfd = accept(int socket, struct sockaddr *address, socklen_t *address_len);
```

### Where:

- `socket`: is the file descriptor of the socket on which to wait for connections (must be bound to address/port).
- `backlog`: maximum length of the queue of pending connections.
- `address`: can be a null pointer, or a pointer to a `sockaddr` structure where the address of the connecting socket shall be returned.
- `address_len`: length of the address.
- `new_sockfd`: new socket descriptor on which the client and the server will communicate.
- `val`: return value which is 0 on success, and -1 otherwise.



# Application Layer

## Sockets: Send and Receive (TCP)

- Processes of sending or receiving messages in TCP are performed by the `send()` and `read()` functions (simpler than UDP ones):

```
#include <sys/socket.h>
```

```
int ob = send(int osock, const void *obuf, size_t olen, int flags);
```

```
#include <unistd.h>
```

```
int ib = read(int isock, void *ibuf, size_t ilen);
```

### Where:

- `osock/isock`: are the file descriptors on which to send/receive a message.
  - `obuf/ibuf`: are pointers to the buffers containing a message to send/receive.
  - `olen/ilen`: are lengths of the messages in bytes.
  - `flags`: specifies the type of message transmission (depends on protocol, typically 0).
  - `ob/ib`: number of bytes that are actually sent/received.
- 
- Notice that we have already specified client and server addresses during the connection phase, we no need to do it here as in `sendto/recvfrom`.

# Application Layer

## TCP Socket Programming (C/C++)

// client-side (includes omitted)

```
int main() {
    int sockfd, status;
    char buffer[1024];
    const char *hello = "Hello from client";
    struct sockaddr_in servaddr;
```

```
    if ( ( sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0 ) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }
```

```
    memset(&servaddr, 0, sizeof(servaddr));
```

```
    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(8080);
    servaddr.sin_addr.s_addr = INADDR_ANY;
```

```
    int n;
```

```
    if ( (status = connect(sockfd, (struct sockaddr*)&servaddr,
        sizeof(servaddr))) < 0 ) {
        printf("\nConnection Failed \n");
        return -1;
    }
```

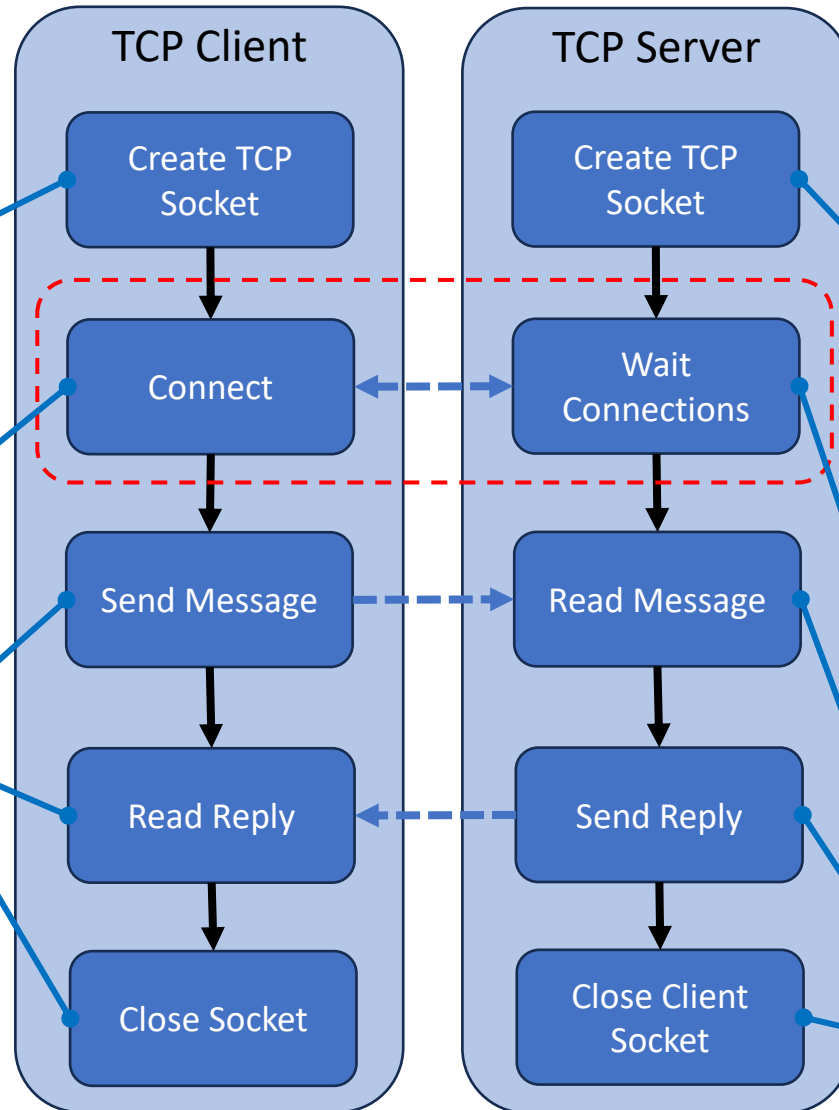
```
    send(sockfd, hello, strlen(hello), 0);
    std::cout<<"Hello message sent."<<std::endl;
```

```
    n = read(sockfd, buffer, 1024);
    std::cout<<"Received \""<<buffer<<"\"<<std::endl;
```

```
    close(sockfd);
```

```
    return 0;
}
```

**Note:** to talk to another computer use `inet_addr()` instead of `INADDR_ANY` (e.g., `inet_addr("192.168.1.10")` )



//server-side (includes omitted)

```
int main() {
    int sockfd, new_socket;
    int opt = 1;
    char buffer[1024];
    const char *hello = "Hello from server";
    struct sockaddr_in servaddr, cliaddr;
```

```
    if ( ( sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0 ) {
        perror("socket creation failed");
        exit(EXIT_FAILURE);
    }
    if (setsockopt(sockfd, SOL_SOCKET, SO_REUSEADDR | SO_REUSEPORT, &opt,
        sizeof(opt))) {
        perror("setsockopt");
        exit(EXIT_FAILURE);
    }
```

```
    memset(&servaddr, 0, sizeof(servaddr));
    memset(&cliaddr, 0, sizeof(cliaddr));
```

```
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = INADDR_ANY;
    servaddr.sin_port = htons(8080);
```

```
    if ( bind(sockfd, (const struct sockaddr *)&servaddr, sizeof(servaddr)) < 0 ) {
        perror("bind failed");
        exit(EXIT_FAILURE);
    }
```

```
    if (listen(sockfd, 3) < 0) {
        perror("listen");
        exit(EXIT_FAILURE);
    }
    int addrlen = sizeof(cliaddr);
    if ( (new_socket = accept(sockfd, (struct sockaddr*)&cliaddr,
        (socklen_t*)&addrlen)) < 0 ) {
        perror("accept");
        exit(EXIT_FAILURE);
    }
```

```
    int n = read(new_socket, buffer, 1024);
    std::cout<<"Received \""<<buffer<<"\"<<std::endl;
```

```
    send(new_socket, hello, strlen(hello), 0);
    std::cout<<"Hello message sent."<<std::endl;
```

```
    close(new_socket);
    close(sockfd);
    return 0;
}
```