

# Computer Network I

Reti di Calcolatori I

Università di Napoli Federico II – Scuola Politecnica e delle Scienze di Base  
Corso di Laurea in Informatica

Riccardo Caccavale  
([riccardo.caccavale@unina.it](mailto:riccardo.caccavale@unina.it))



# Transport Layer

## Reliable Data Transfer Problem

- Let's assume we are at the railway station waiting for the train number 6 which should arrive on platform 5.
- The train changes platform, it will arrive on platform 9 instead of platform 5, a message is sent from station control to communicate it:

“Train 6 will arrive on platform 9”

- What happens if the communication is unreliable:
  - Words could be lost: “Train %&! will arrive on platform 9”
  - Words could be altered: “Train 7 will arrive on platform 9”
  - Words could be swapped: “Train 9 will arrive on platform 7”
  - Words could be duplicated: “Train 66 will arrive on platform 9”
- It is possible you can understand that the message is wrong (for example in case 1), but you may also take the wrong train or miss the right one.



# Transport Layer

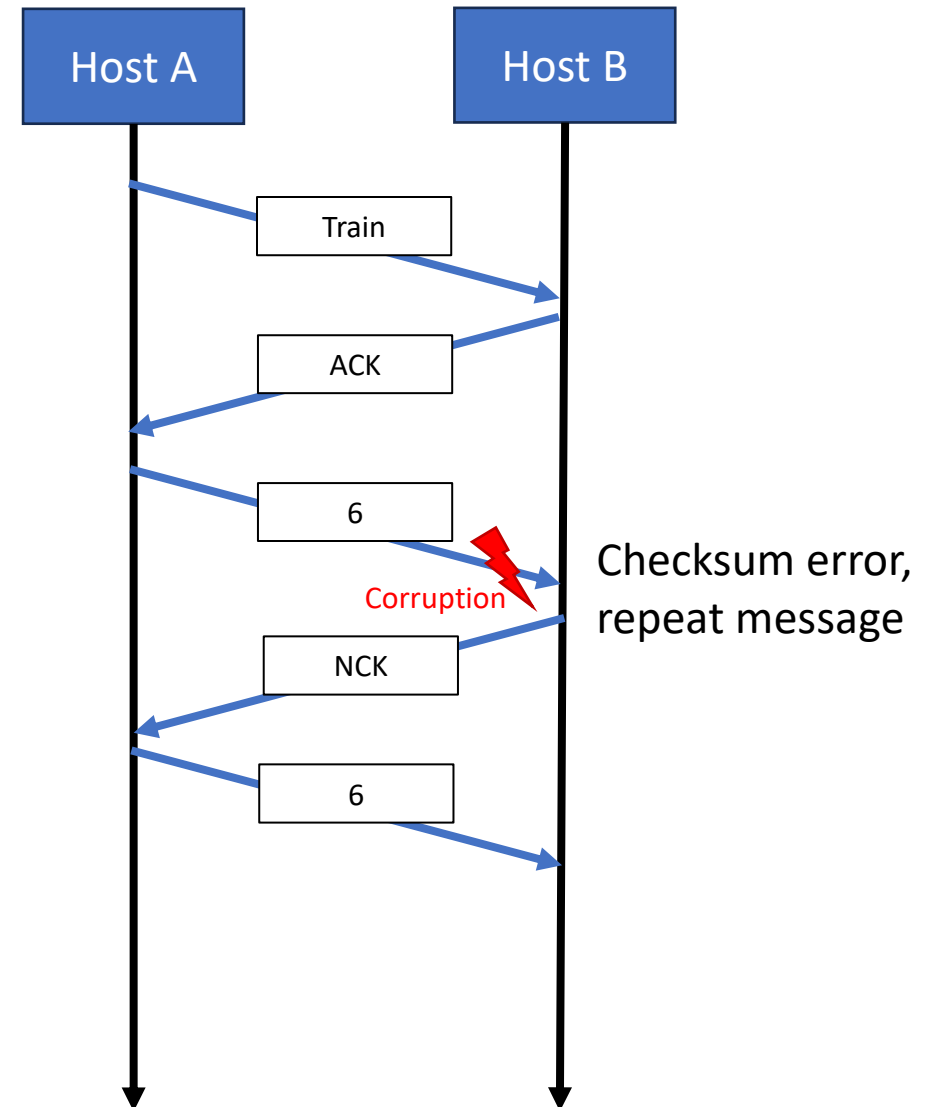
## Reliable Data Transfer

- Error control (e.g., through checksum) allows receiver at least to understand **if the message is corrupted**, but this is not enough. Reliable data transfer **is still one of the main problem of networking**.
  - The problem is not only addressed at the transport layer, but also at the link layer and (often) at the application layer.
- In a reliable channel 3 elements must be guaranteed:
  1. None of the transmitted bits **are corrupted** (flipped from 0 to 1, or vice versa).
  2. None of the transmitted bits **are lost or repeated**.
  3. All bits are delivered **in the exact order** in which they were sent.
- In general, we must assume the lower-level **network layer to be unreliable**.
  - This is a realistic assumption, for example, TCP is a reliable data transfer protocol that **is implemented on top of an unreliable (IP) end-to-end network layer**.

# Transport Layer

## Reliable Data Transfer: Packet Corruption and Stop-and-wait

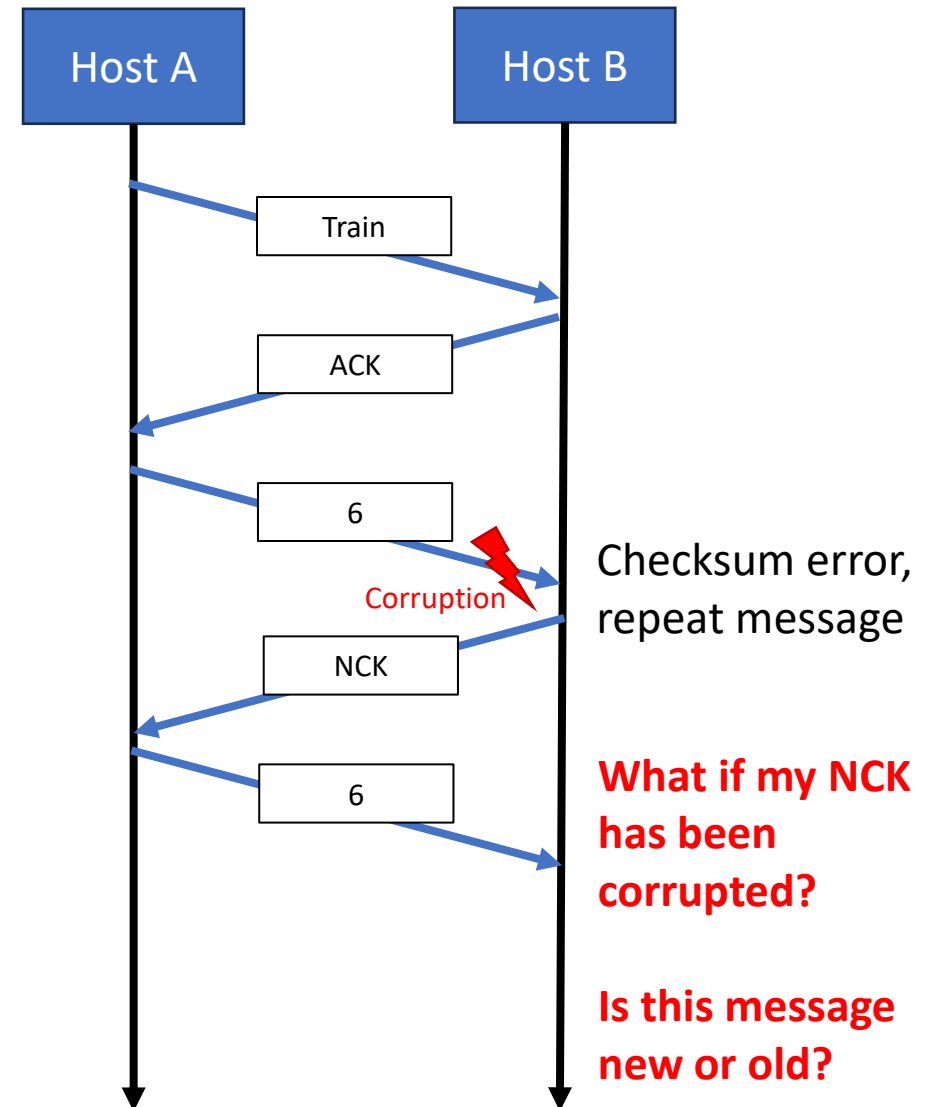
- Let's assume we have a **channel where bits may only be corrupted** (not lost).
  - Corruption of bits is a typical problem, which is **mainly due to the physical components** of a network as a packet is transmitted, propagated, or buffered several time during the communication.
- **Stop-and-wait:** a first approach requires each message to be acknowledged before sending a new one:
  - **Positive acknowledgments** (ACK) means that the message has been received intact.
  - **Negative acknowledgments** (NCK) means that an error occurred, so the message must be repeated.
- In a computer network, **protocols based on retransmission are also known as ARQ** (Automatic Repeat reQuest) protocols.



# Transport Layer

## Reliable Data Transfer: Stop-and-wait

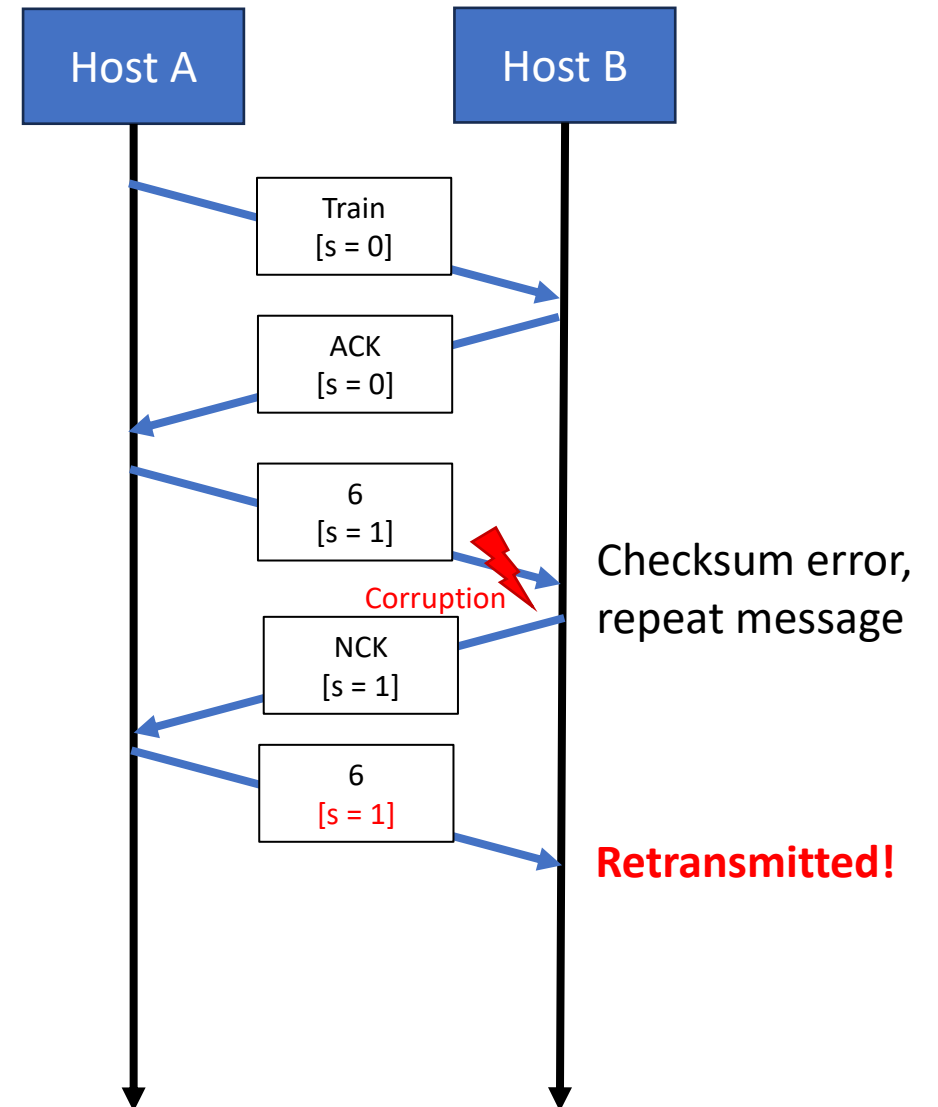
- There are a couple of **issues** with this approach:
  - What happens if the **ACK/NCK message is corrupted** itself?
  - How can Host B be sure that a **message is the repetition instead of a new message**?
- **Sequence number**: add a new field to the headers of packets which specifies the order in which packages should be received.



# Transport Layer

## Reliable Data Transfer: Stop-and-wait

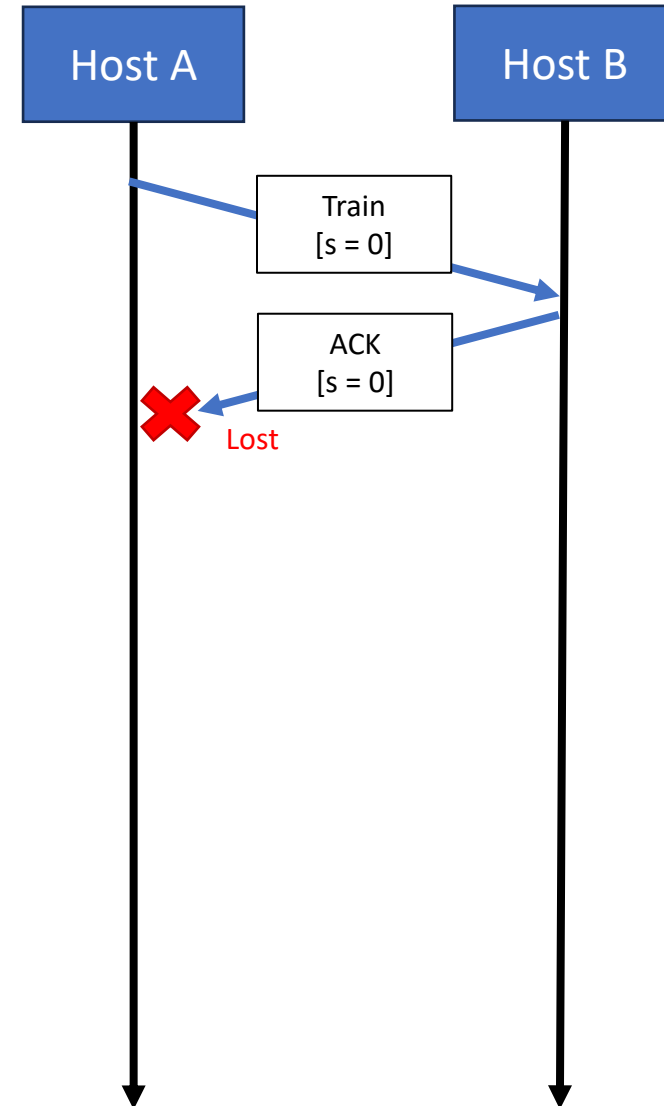
- In a stop-and wait approach we always have **one message** on the line.
- There is no need to specify the whole sequence of packets, we just need to **differentiate between the current packet and the previous one**.
- Therefore, in this case, **we just need to add one bit** ( $s = 0/1$ ) to the headers of packets.



# Transport Layer

## Reliable Data Transfer: Packet Loss and Stop-and-wait

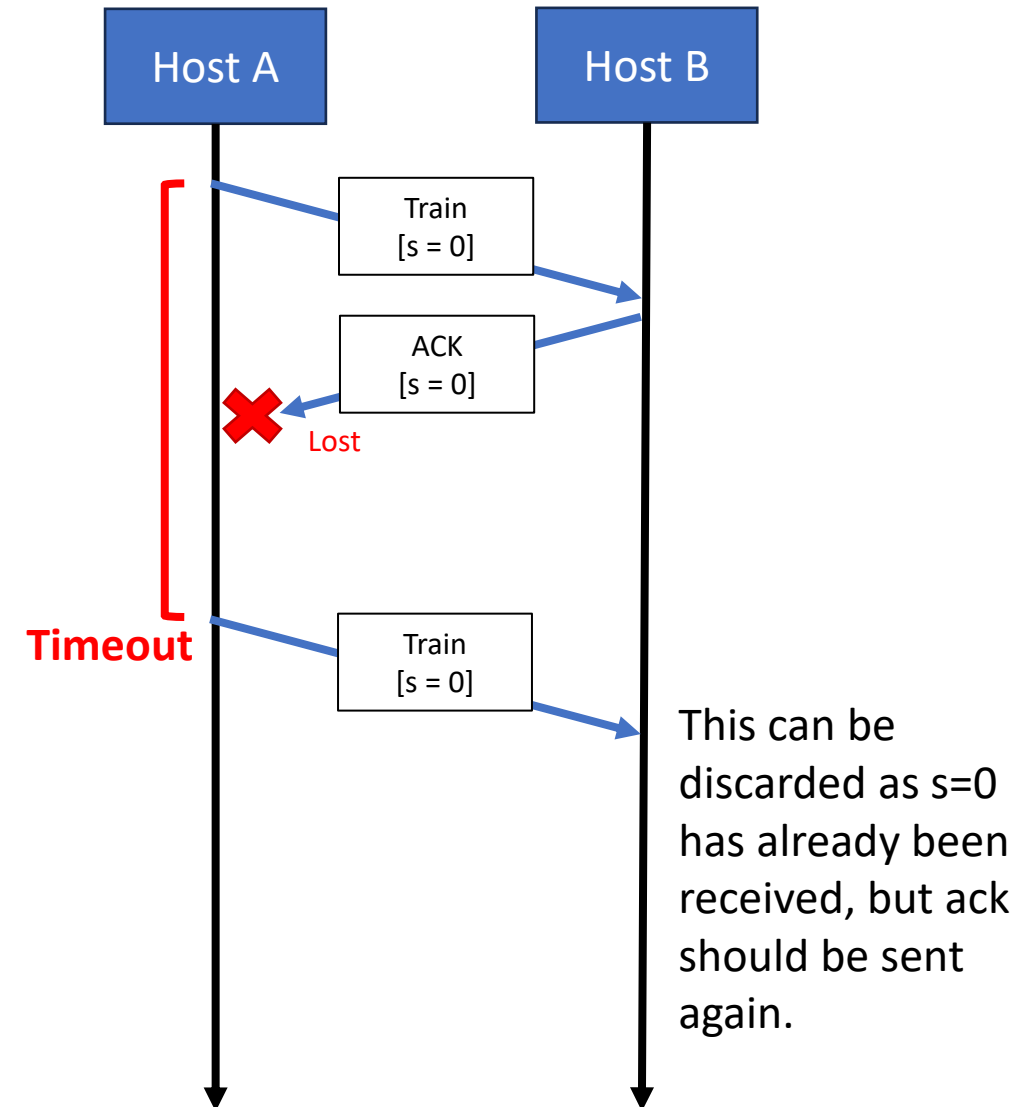
- Let's assume we now have a **channel where packets may also be lost**.
  - This is also quite reasonable as **network devices may have buffer overflow** in case of intense traffic.
- Here we have a clear problem with stop-and-wait approach: the loop. **If one message is lost, hosts will never send a new one**.
  - Notice that this happens whether we loose message or acknowledgment, as in both cases host A will not be triggered to send a new message.



# Transport Layer

## Reliable Data Transfer: Packet Loss and Stop-and-wait

- **Timeout:** a very simple yet effective solution is to **add a timeout on the sender-side** that, once expired, allows the Host to try again.
  - A similar approach can be used in UDP-based DNS.
- This works **both in the case of packet loss and in acknowledgment loss.**
  - In case of duplicate the receiver has just to discard the packet.

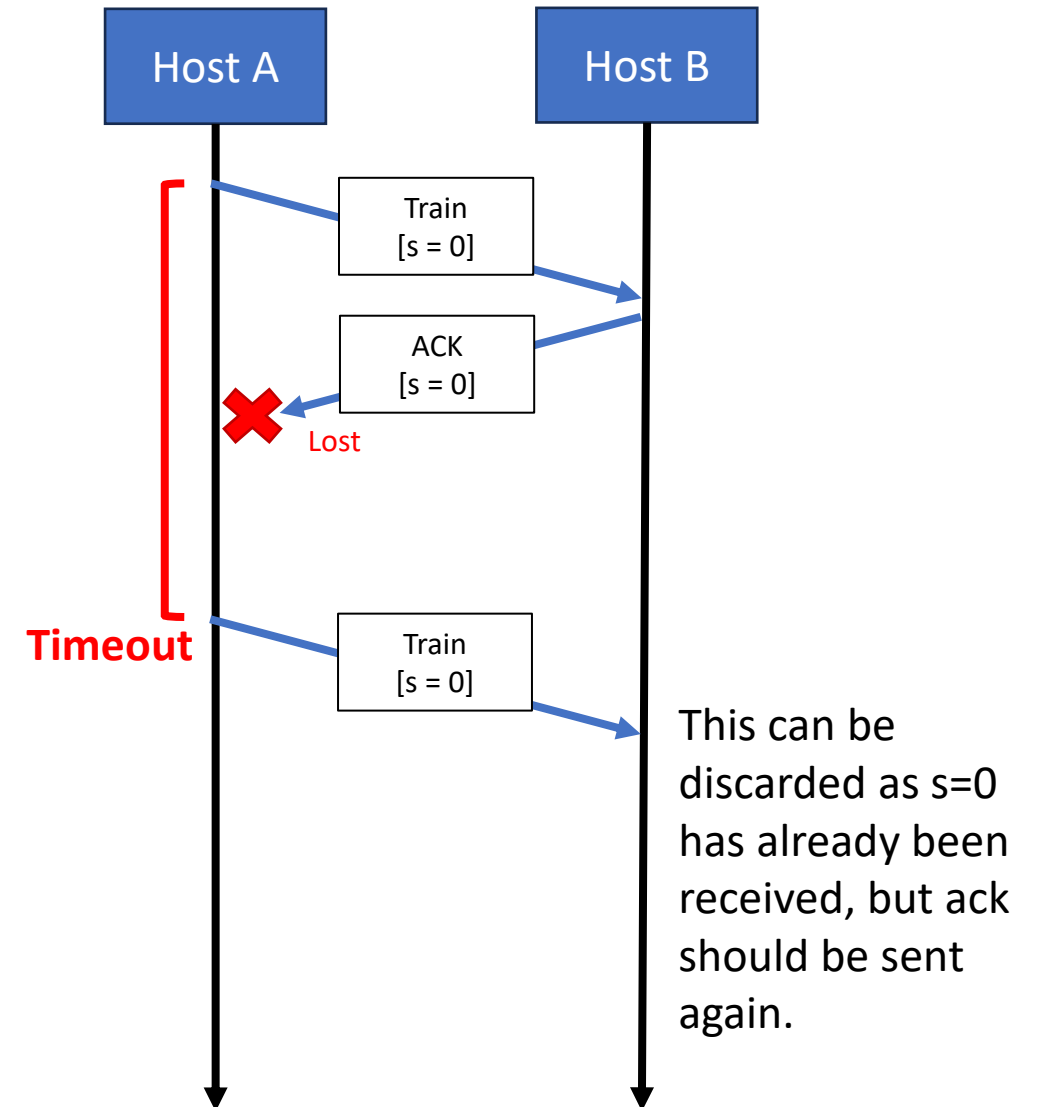




# Transport Layer

## Reliable Data Transfer: Packet Loss and Stop-and-wait


- The challenge here is **how to estimate a suitable timeout** (it depends on the RTT).
  - Too long timeout causes communication to be slowed.
  - Too short timeout causes the packets to overlap.
- We can say that a reasonable timeout should be somehow longer than the RTT.
- Timeout saves stop-and-wait from loop, but the **performance are quite bad...**



# Transport Layer

## Reliable Data Transfer: Water Example


- We can see **data transfer as a water flow problem**:
  - There is **the time ( $t$ ) to move the water from the tank to the tube**. While entering the water is also moving toward the tube.
  - After  $t$  **all the water from the tank is inside the tube** (the initial tank is empty as well as the destination tank).
  - It **takes a certain time ( $RTT/2$ ) for the water to flow** into the tube.
  - Finally, **the water arrives at the destination tank** and (by assuming same download/upload time) it takes the same time  $t$  to pour out of the tube ( $RTT/2 + t$  to receive the last drop).



# Transport Layer

## Reliable Data Transfer: Performance of Stop-and-Wait

- Let's consider the idealized case of **two hosts located on the opposite coasts** of the United States.




- The speed-of-light RTT between these two end systems is **approximately 30 milliseconds** (0.03 sec.). Let's assume to have:
  - A channel with a transmission rate ( $R$ ) of 1 Gbps ( $10^9$  bits per second).
  - A packet size ( $L$ ) of 1000 bytes (8000 bits).
- The time ( $t$ ) needed to transmit the packet to the channel is:

$$t = \frac{L}{R} = \frac{8000}{10^9} = 0.000008 \text{ (8 microseconds)}$$

# Transport Layer

## Reliable Data Transfer: Performance of Stop-and-Wait


- Let's consider the idealized case of **two hosts located on the opposite coasts of the United States**.
- In a **stop-and-wait** protocol:
  - the **sender begins sending** the packet at  $t_0$ ,
  - the **last bit enters the channel** at  $t_0 + 0.000008$  sec.,
  - the **packet takes 0.015 sec. to reach** the receiver,
  - the **last bit is received** at  $t = \text{RTT}/2 + L/R = 0.015008$  sec.,
  - assuming that ACK packets are extremely small (their **transmission time can be neglected**) the ACK emerges back at the sender at  $t = \text{RTT} + L/R = 0.030008$  sec.
- In 0.030008 sec. of total transmission time, the sender was waiting almost all the time (99.973% of the time).



# Transport Layer

## Reliable Data Transfer: Performance of Pipelining


- **Pipelining:** instead of stop-and-wait, the sender is **allowed to send multiple packets** without waiting for acknowledgments.
- In pipelining approaches (e.g., 3 packets instead of 1):
  - the **sender begins sending the 3 packets** at  $t_0$ ,
  - the **last bit of the last packet enters the channel** at  $t_0 + 0.000024$  sec.,
  - the **packets take 0.015 sec. to reach** the receiver,
  - the **last bit is received** at  $t = \text{RTT}/2 + L/R = 0.015024$  sec.,
  - assuming that ACK packets are extremely small (their **transmission time can be neglected**) the ACK emerges back at the sender at  $t = \text{RTT} + L/R = 0.030024$  sec.
- In 0.030024 sec. **the sender was waiting 0.035% less** (99.920% instead of 99.973%).



# Transport Layer

## Reliable Data Transfer: Performance of Pipelining


- **Pipelining:** instead of stop-and-wait, the sender is **allowed to send multiple packets** without waiting for acknowledgments.
- In pipelining approaches:
  - The **range of sequence numbers must be increased** since each in-transit packet (not counting retransmissions) must have a unique sequence number and there may be multiple, in-transit, unacknowledged packets.
  - We **need buffers to store incoming packets** since we don't know if packets are correct or there are "holes" in the transmission.
- There are 2 basic protocols using pipelining:
  - **Go-Back-N.**
  - **Selective Repeat.**



# Transport Layer

## Reliable Data Transfer: Go-Back-N

- **Go-Back-N (GBN)** protocol (aka sliding-window protocol): the sender is allowed to transmit multiple packets (when available) **without waiting for an acknowledgment** but is constrained to have **no more than N unacknowledged packets**.

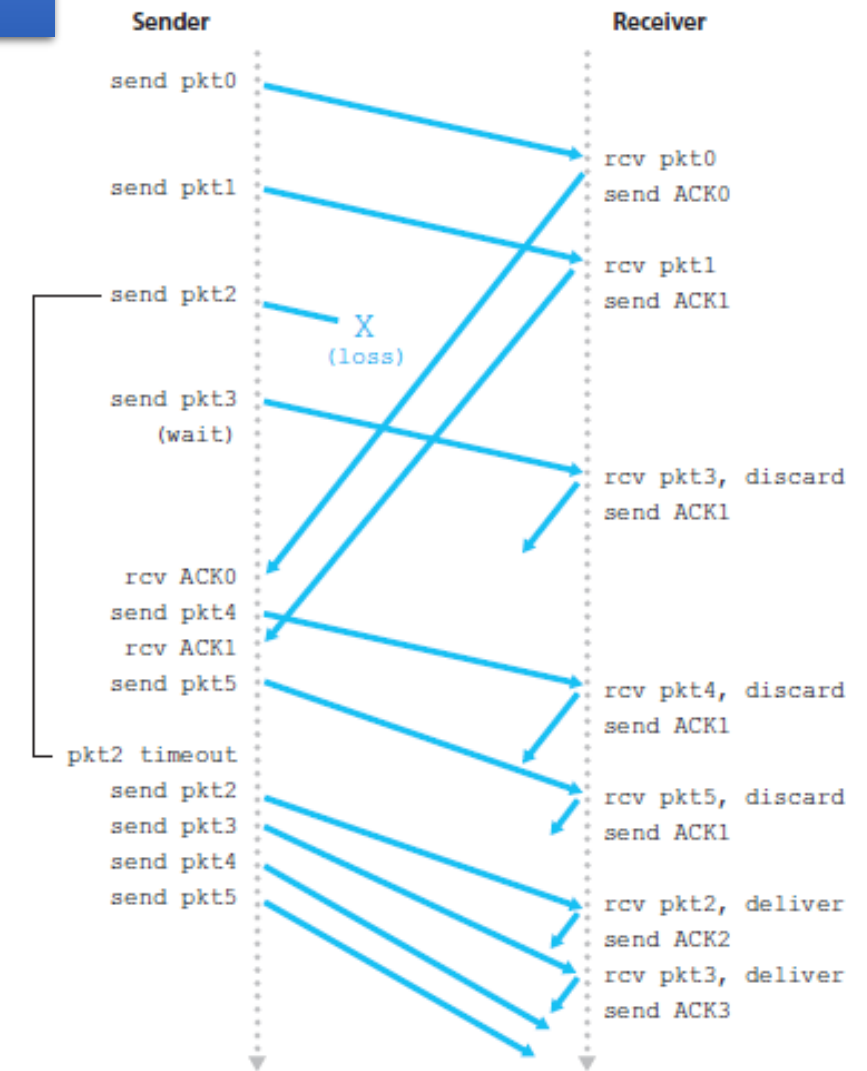


- Base: is the **sequence number of the oldest unacknowledged packet**.
- Nextseqnum: is the **smallest unused sequence number** (next to be sent).
- We have that:
  - Packets in  $[0, \text{base}-1]$  have been **transmitted and acknowledged**.
  - Packets in  $[\text{base}, \text{nextseqnum}-1]$  have been **sent but not yet acknowledged**.
  - Packets in  $[\text{nextseqnum}, \text{base}+N-1]$  **can be sent**.
  - Packets in  $[\text{base}+N, +\infty]$  **cannot be used** until a new acknowledgment is received.

# Transport Layer

## Reliable Data Transfer: Go-Back-N

- What about the receiver? In GBN protocol **receiver is quite simple**.
- The receiver has **simply to discard out-of-order packets** (whether they are damaged or not) and to deliver to the upper level (application) in-order packets only.
  - Discarded packets (not acknowledged) will be eventually retransmitted by the sender.
- With this approach good packets are also discarded but the overall process is quite simple:
  - the **sender must maintain the indices of the window**,
  - the **receiver only needs to maintain the sequence number** of the next in-order packet.



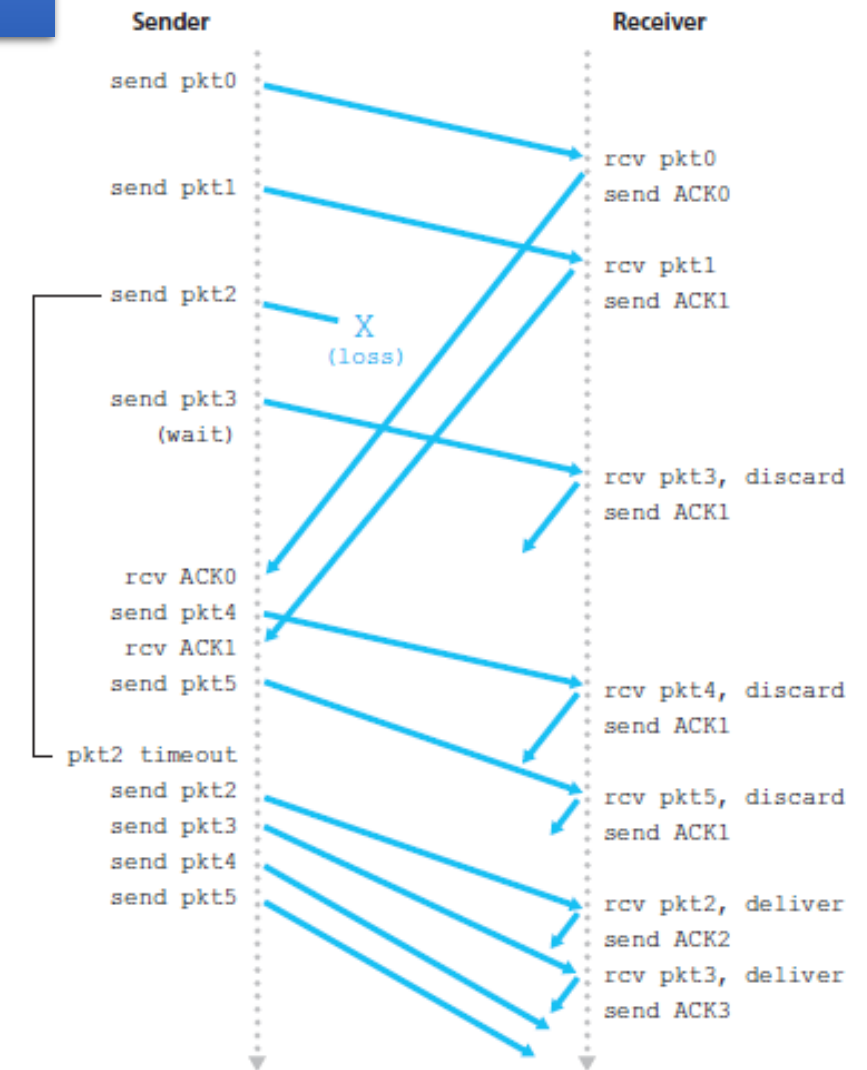
Since pkt2 is lost, all pkt3, pkt4, pkt5, must be discarded even if correct.



# Transport Layer

## Reliable Data Transfer: Go-Back-N

- Of course, the disadvantage of throwing away a correctly received packet is that **we need to resend it again**.
- This is a **chain-reaction**, if we are losing packets due to network difficulties many correctly received but out-of-order packets may be discarded forcing the sender to retransmit them.
- The retransmission itself might be lost or damaged and thus **even more retransmissions would be required**.

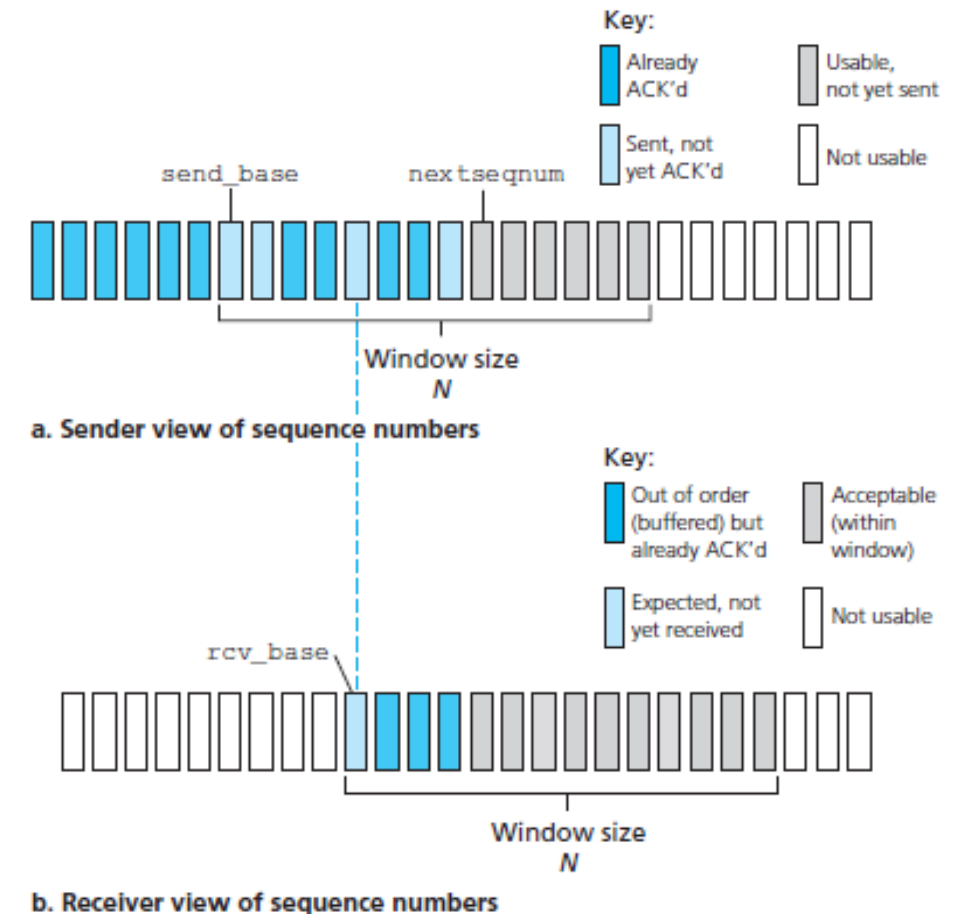


Not only pkt2, but also pkt3, pkt4, pkt5, must be transmitted again.

# Transport Layer

## Reliable Data Transfer: Selective Repeat

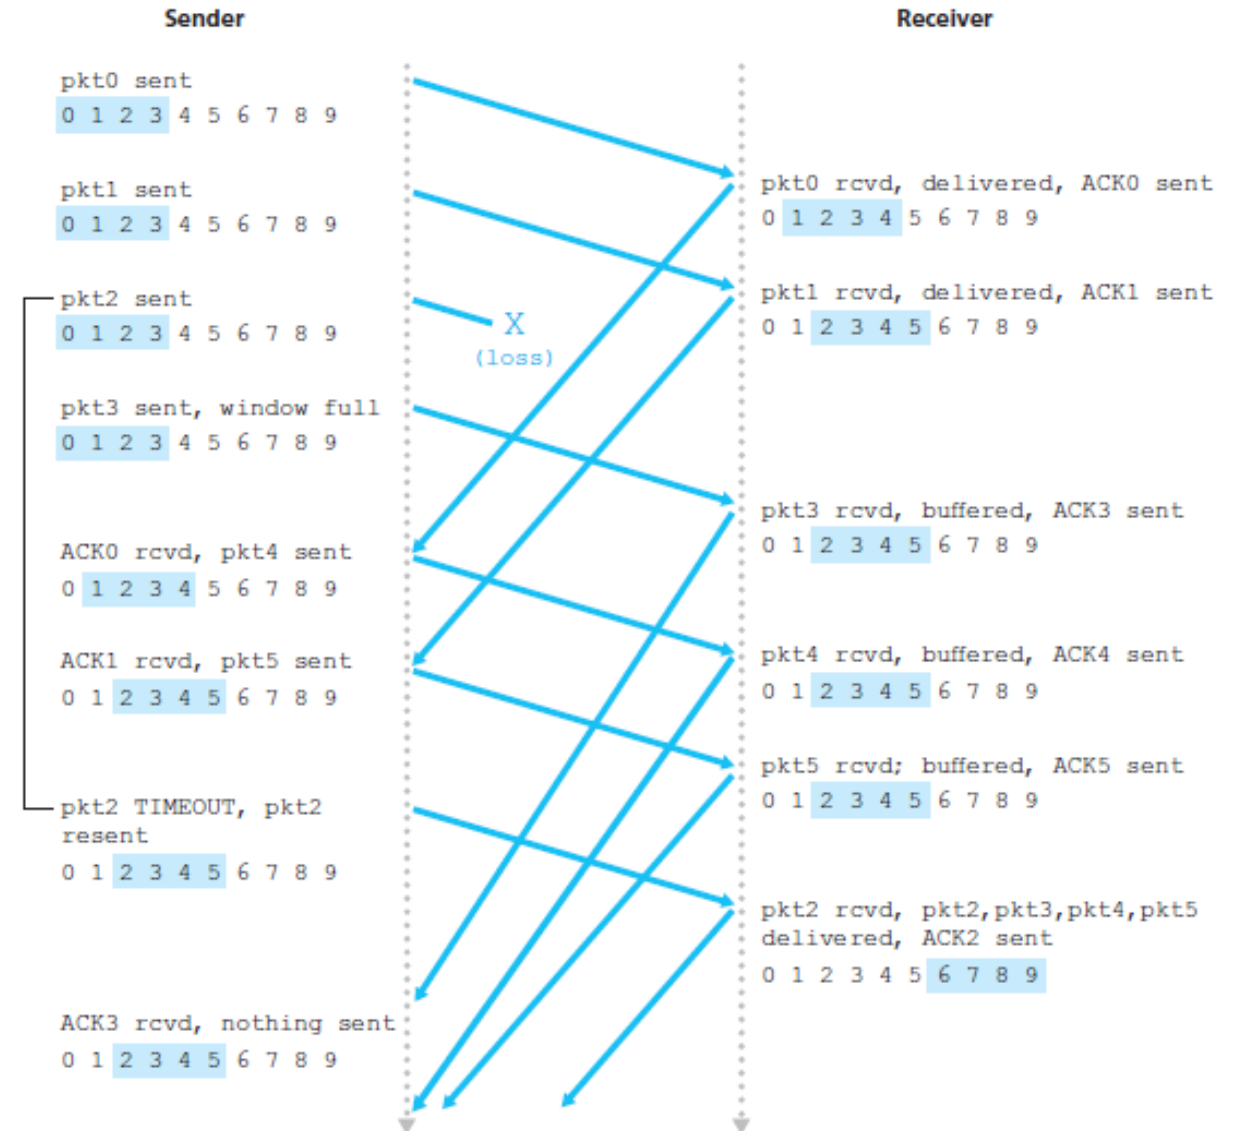
- Despite GBN is more effective than stop-and-wait, it may **suffer from performance problems**.
- Considering large window size, a **single packet error can cause GBN to retransmit a large number of packets**, many unnecessarily.
- In **selective-repeat (SR)** protocols: the sender retransmit **only those packets that were likely received in error** (lost or corrupted):
  - As in GBN **individual packets are acknowledged**.
  - A **window size of N** is again used to limit the number of unacknowledged packets.
  - Unlike GBN, **out-of-order packets are buffered** (stored) and acknowledged by the receiver.



# Transport Layer

## Reliable Data Transfer: Selective Repeat

- Example of SR operation in the presence of lost packets: the receiver initially buffers pkt3, pkt4, and pkt5, while waiting for pkt2 (lost) to be retransmitted.
- Despite the previous example, here we avoid to resend pkt3, pkt4, pkt5, in so avoiding additional transmission loss or errors.



# Transport Layer

## Reliable Data Transfer: Selective Repeat

- One issue in SR is that the window size is related to the sequence number, **and sequence number is finite.**
- In this example we have four packets, **a max sequence number of 3** and a window size of 3.
- Let's assume packets 0 to 2 are correctly received and acknowledged, hence receiver's window moves to 6<sup>th</sup> packet (i.e., to packets 3, 0, 1):
  - Case a:** the ACKs for the first three packets are lost and the sender retransmits these packets. *Is packet 0 new or old?*
  - Case b:** the ACKs are received, packets 3 and 0 are sent but packet 3 is lost. *Is packet 0 new or old?*
- To avoid this issue, **the sequence number must be at least 2 times the window size.**

